

## learningAgents.py ([original](#))

---

```
# learningAgents.py
# -----
# Licensing Information: Please do not distribute or publish solutions to this
# project. You are free to use and extend these projects for educational
# purposes. The Pacman AI projects were developed at UC Berkeley, primarily by
# John DeNero (denero@cs.berkeley.edu) and Dan Klein (klein@cs.berkeley.edu).
# For more info, see http://inst.eecs.berkeley.edu/~cs188/sp09/pacman.html

from game import Directions, Agent, Actions

import random,util,time

class ValueEstimationAgent(Agent):
    """
    Abstract agent which assigns values to (state,action)
    Q-Values for an environment. As well as a value to a
    state and a policy given respectively by,

    
$$V(s) = \max_{a \in \text{actions}} Q(s,a)$$

    
$$\text{policy}(s) = \arg\max_{a \in \text{actions}} Q(s,a)$$


    Both ValueIterationAgent and QLearningAgent inherit
    from this agent. While a ValueIterationAgent has
    a model of the environment via a MarkovDecisionProcess
    (see mdp.py) that is used to estimate Q-Values before
    ever actually acting, the QLearningAgent estimates
    Q-Values while acting in the environment.
    """

    def __init__(self, alpha=0.5, epsilon=0.05, gamma=0.9, numTraining = 10):
        """
        Sets options, which can be passed in via the Pacman command line using -a
        alpha=0.5,...
        alpha      - learning rate
        epsilon    - exploration rate
        gamma      - discount factor
        numTraining - number of training episodes, i.e. no learning after these many
        episodes
        """
        self.alpha = float(alpha)
        self.epsilon = float(epsilon)
        self.gamma = float(gamma)
        self.numTraining = int(numTraining)

    #####
    # Override These Functions #
    #####
    def getQValue(self, state, action):
        """
        Should return Q(state,action)
        """
        util.raiseNotDefined()

    def getValue(self, state):
        """
        What is the value of this state under the best action?
        Concretely, this is given by

        
$$V(s) = \max_{a \in \text{actions}} Q(s,a)$$

        """
        util.raiseNotDefined()

    def getPolicy(self, state):
        """
```

What is the best action to take in the state. Note that because we might want to explore, this might not coincide with `getAction`. Concretely, this is given by

$$\text{policy}(s) = \arg\max_{\{a \text{ in actions}\}} Q(s,a)$$

If many actions achieve the maximal Q-value, it doesn't matter which is selected.

"""

`util.raiseNotDefined()`

**def** `getAction(self, state):`

"""

state: can call `state.getLegalActions()`

Choose an action and return it.

"""

`util.raiseNotDefined()`

**class** `ReinforcementAgent(ValueEstimationAgent):`

"""

Abstract Reinforcement Agent: A `ValueEstimationAgent`

which estimates Q-Values (as well as policies) from experience rather than a model

What you need to know:

- The environment will call

`observeTransition(state, action, nextState, deltaReward)`,

which will call `update(state, action, nextState, deltaReward)`

which you should override.

- Use `self.getLegalActions(state)` to know which actions are available in a state

"""

#####

*# Override These Functions*

#####

**def** `update(self, state, action, nextState, reward):`

"""

This class will call this function, which you write, after observing a transition and reward

"""

`util.raiseNotDefined()`

#####

*# Read These Functions*

#####

**def** `getLegalActions(self, state):`

"""

Get the actions available for a given state. This is what you should use to obtain legal actions for a state

"""

**return** `self.actionFn(state)`

**def** `observeTransition(self, state, action, nextState, deltaReward):`

"""

Called by environment to inform agent that a transition has been observed. This will result in a call to `self.update` on the same arguments

NOTE: Do *\*not\** override or call this function

"""

`self.episodeRewards += deltaReward`

`self.update(state, action, nextState, deltaReward)`

**def** `startEpisode(self):`

"""

Called by environment when new episode is starting

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

```

    """
    self.lastState = None
    self.lastAction = None
    self.episodeRewards = 0.0

def stopEpisode(self):
    """
    Called by environment when episode is done
    """
    if self.episodesSoFar < self.numTraining:
        self.accumTrainRewards += self.episodeRewards
    else:
        self.accumTestRewards += self.episodeRewards
    self.episodesSoFar += 1
    if self.episodesSoFar >= self.numTraining:
        # Take off the training wheels
        self.epsilon = 0.0 # no exploration
        self.alpha = 0.0 # no learning

def isInTraining(self):
    return self.episodesSoFar < self.numTraining

def isInTesting(self):
    return not self.isInTraining()

def __init__(self, actionFn = None, numTraining=100, epsilon=0.5, alpha=0.5,
gamma=1):
    """
    actionFn: Function which takes a state and returns the list of legal actions
    alpha - learning rate
    epsilon - exploration rate
    gamma - discount factor
    numTraining - number of training episodes, i.e. no learning after these many
    episodes
    """
    if actionFn == None:
        actionFn = lambda state: state.getLegalActions()
    self.actionFn = actionFn
    self.episodesSoFar = 0
    self.accumTrainRewards = 0.0
    self.accumTestRewards = 0.0
    self.numTraining = int(numTraining)
    self.epsilon = float(epsilon)
    self.alpha = float(alpha)
    self.gamma = float(gamma)

#####
# Controls needed for Crawler #
#####
def setEpsilon(self, epsilon):
    self.epsilon = epsilon

def setLearningRate(self, alpha):
    self.alpha = alpha

def setDiscount(self, discount):
    self.gamma = discount

def doAction(self, state, action):
    """
    Called by inherited class when
    an action is taken in a state
    """
    self.lastState = state
    self.lastAction = action

#####
# Pacman Specific #

```

```
#####
def observationFunction(self, state):
    """
    This is where we ended up after our last action.
    The simulation should somehow ensure this is called
    """
    if not self.lastState is None:
        reward = state.getScore() - self.lastState.getScore()
        self.observeTransition(self.lastState, self.lastAction, state, reward)
    return state

def registerInitialState(self, state):
    self.startEpisode()
    if self.episodesSoFar == 0:
        print 'Beginning %d episodes of Training' % (self.numTraining)

def final(self, state):
    """
    Called by Pacman game at the terminal state
    """
    deltaReward = state.getScore() - self.lastState.getScore()
    self.observeTransition(self.lastState, self.lastAction, state, deltaReward)
    self.stopEpisode()

    # Make sure we have this var
    if not 'episodeStartTime' in self.__dict__:
        self.episodeStartTime = time.time()
    if not 'lastWindowAccumRewards' in self.__dict__:
        self.lastWindowAccumRewards = 0.0
    self.lastWindowAccumRewards += state.getScore()

NUM_EPS_UPDATE = 100
if self.episodesSoFar % NUM_EPS_UPDATE == 0:
    print 'Reinforcement Learning Status:'
    windowAvg = self.lastWindowAccumRewards / float(NUM_EPS_UPDATE)
    if self.episodesSoFar <= self.numTraining:
        trainAvg = self.accumTrainRewards / float(self.episodesSoFar)
        print '\tCompleted %d out of %d training episodes' % (
            self.episodesSoFar, self.numTraining)
        print '\tAverage Rewards over all training: %.2f' % (
            trainAvg)
    else:
        testAvg = float(self.accumTestRewards) / (self.episodesSoFar -
self.numTraining)
        print '\tCompleted %d test episodes' % (self.episodesSoFar -
self.numTraining)
        print '\tAverage Rewards over testing: %.2f' % testAvg
        print '\tAverage Rewards for last %d episodes: %.2f' % (
            NUM_EPS_UPDATE, windowAvg)
        print '\tEpisode took %.2f seconds' % (time.time() - self.episodeStartTime)
        self.lastWindowAccumRewards = 0.0
        self.episodeStartTime = time.time()

if self.episodesSoFar == self.numTraining:
    msg = 'Training Done (turning off epsilon and alpha)'
    print '%s\n%s' % (msg, '-' * len(msg))
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder